Extended abstract

Simon Théorêt

Abstract

Data augmentation is a key component to training robust models and to prevent overfitting in computer vision. The Fourier perspective [Yin et al., 2019] gave us a better understanding of the processs behind such improvements by showing us the tradeoffs between protecting against corruptions in low frequency and high frequency and the limitations of data augmentation. These insights also raises questions about the bias of gradients towards low frequencies: Does the optimizer and the architecture of a model bias the model to rely on low frequency features. In this work, we investigate the impact of Gaussian data augmentation and adversarial training on a different set of architectures and optmizers. We provide an hypothesis that gradients of adversarially trained models and models trained on Gaussian augmented data are naturally biased towards low frequency features, as they contain more relevant information for classification. To test our hypothesis, we provide an experimental protocol for testing our hypothesis against different architectures and optmizers by computing the accuracy of the trained models on images containing either high frequency features or low frequency features.

1   Introduction

The fourrier perspective introduced by Yin et al. [2019] paved the way for exploring the Fourier space of images' features and how these features are used by models. This approach was used to classify features in two categories: high frequency and low frequency features. The first category includes features such as images' texture, and the second is related to the countours and shapes in images, as stated in Krishnamachari et al. [2023]. Although high frequency features are often invisible to the human eye, the Fourier perspective showed that these features could be successfully used by Convolution Neural Networks (CNN) in image classification. However, high frequency features are not robust, as showed by Zhang and Zhu [2019] and yet models are often biased toward using these features. On the contrary, low frequency features, such as shape, are often the preferred features of adversarially trained neural networks or network trained with a Gaussian augmented dataset. The Fourier perspective article Yin et al. [2019] limited their experiments to the ResNet architecture. The lack of empirical research using different architectures and optmizers prompts the question:

Does the architecture and the optimizer influence the bias toward low frequency features in adversarially trained models and models trained on Gaussian augmented datasets?

2   Definitions

We use the following notations: $\mathcal{F} : \mathbb{R}^{d_1 \times d_2} \to \mathbb{C}^{d_1 \times d_2}$ denotes the discrete fourrier transform (DFT) of an image and we omit the dimensions of the channels, as every channel is treated independently of the other channels. When we visualize the Fourier spectrum, we always shift the low frequency components to the center of the spectrum and only show the magnitude of the images, not the phase.

To filter components of an image based on their frequency, we use the methodology of Yin et al. [2019]: We set to 0 every point in the Fourier spectrum that is not in the square of width B centered at the highest/lowest frequency. We then apply the inverse DFT to recover the original image, with the low/high frequency components filtered out.

Our Gaussian augmentation method follows the methodology of Yin et al. [2019]. We assume that pixels take values in the range $[0, 1]$. Pixel values are always clipped to that range. Gaussian data

augmentation with parameters σ is defined as the following operation: i.i.d. Gaussian noise $\mathcal{N}(0, \tilde{\sigma}^2)$, is applied at each iteration and at every pixel. The value of $\tilde{\sigma}^2$ is chosen uniformly at random from $[0, \sigma]$.

3 Problem statement and related works

Our goal is to assert whether or not optimizers and architecture influence the bias in the features selection of models trained with a Gaussian augmented dataset or adversarially trained. Precedent works have tried to formalize the Fourier sensitivity of CNN with the works of Krishnamachari et al. [2023] and gave experimental insights regarding the tunining of models towards certain frequencies, thanks to the works of Krishnamachari et al. [2023], Geirhos et al. [2022], Yin et al. [2019], Mo et al. [2022]. Notably, the works of Park and Kim [2022] demonstrated that the multi-head self-attentions, such as ViT reduce high frquency signals, while CNN amplify them.

4 Proposed experimental protocol

We are planning on using two different architectures for our experiments: the ALL-CNN architecture Springenberg et al. [2015], which is mostly convolution layers stacked on top of each others, and the mobileViT architecture, which introduce some transformer modules into the architecture Mehta and Rastegari [2022]. We do not include any pure ViT (vision transformers) models, due to their heavy computational demands.

Our experiments consists of training a total of 12 models on the MNIST dataset and evaluating the performance on the validation set with only high frequency features or low frenquecy features. For both the CNN architecture and the ViT architecture, each are trained on the default MNIST dataset, adversarially trained on the MNIST dataset and trained on the MNIST dataset with Gaussian augmentation. This procedure is repeated for two Stochastic gradient descent with momentum and AdamW. Both of these optimizers were used in the original training of ALL-CNN and MobileViT models, respectively.

To accomplish these experiments, we have access to Google Colab Pro account and a NVDIA GTX 3060 GPU.

References

R. Geirhos, P. Rubisch, C. Michaelis, M. Bethge, F. A. Wichmann, and W. Brendel. Imagenet-trained cnns are biased towards texture; increasing shape bias improves accuracy and robustness, 2022.

K. Krishnamachari, S.-K. Ng, and C.-S. Foo. Fourier sensitivity and regularization of computer vision models, 2023.

S. Mehta and M. Rastegari. Mobilevit: Light-weight, general-purpose, and mobile-friendly vision transformer, 2022.

Y. Mo, D. Wu, Y. Wang, Y. Guo, and Y. Wang. When adversarial training meets vision transformers: Recipes from training to architecture, 2022.

N. Park and S. Kim. How do vision transformers work?, 2022.

J. T. Springenberg, A. Dosovitskiy, T. Brox, and M. Riedmiller. Striving for simplicity: The all convolutional net, 2015.

D. Yin, R. G. Lopes, J. Shlens, E. D. Cubuk, and J. Gilmer. A fourier perspective on model robustness in computer vision. *CoRR*, abs/1906.08988, 2019. URL http://arxiv.org/abs/1906.08988.

T. Zhang and Z. Zhu. Interpreting adversarially trained convolutional neural networks, 2019.